

Lecture 03: Dynamic Programming

Davit Ghazaryan

February 11, 2025



Table of contents

- 1 Introduction
- 2 Policy evaluation (Prediction)
- 3 Policy improvement (Control)
- 4 Policy and value iteration
- 5 Further aspects

What is dynamic programming (DP)?

Basic DP definition

- ▶ **Dynamic**: sequential or temporal problem structure
- ▶ **Programming**: mathematical optimization, i.e., numerical solutions

What is dynamic programming (DP)?

Basic DP definition

- ▶ **Dynamic**: sequential or temporal problem structure
- ▶ **Programming**: mathematical optimization, i.e., numerical solutions

Further characteristics:

- ▶ DP is a collection of algorithms to solve MDPs and neighboring problems.
 - ▶ **We will focus only on finite MDPs.**
 - ▶ In case of continuous action/state space: apply quantization.

What is dynamic programming (DP)?

Basic DP definition

- ▶ **Dynamic**: sequential or temporal problem structure
- ▶ **Programming**: mathematical optimization, i.e., numerical solutions

Further characteristics:

- ▶ DP is a collection of algorithms to solve MDPs and neighboring problems.
 - ▶ We will focus only on finite MDPs.
 - ▶ In case of continuous action/state space: apply quantization.
- ▶ Use of value functions to organize and structure the search for an optimal policy.
- ▶ Breaks problems into subproblems and solves them.

Requirements for DP

DP can be applied to problems with the following characteristics.

- ▶ Optimal substructure:
 - ▶ Principle of optimality applies.
 - ▶ Optimal solution can be derived from subproblems.

Requirements for DP

DP can be applied to problems with the following characteristics.

- ▶ Optimal substructure:
 - ▶ Principle of optimality applies.
 - ▶ Optimal solution can be derived from subproblems.
- ▶ Overlapping subproblems:
 - ▶ Subproblems recur many times.
 - ▶ Hence, solutions can be cached and reused.

Requirements for DP

DP can be applied to problems with the following characteristics.

- ▶ Optimal substructure:
 - ▶ Principle of optimality applies.
 - ▶ Optimal solution can be derived from subproblems.
- ▶ Overlapping subproblems:
 - ▶ Subproblems recur many times.
 - ▶ Hence, solutions can be cached and reused.

How is that connected to MDPs?

- ▶ MDPs satisfy above's properties:
 - ▶ Bellman equation provides recursive decomposition.
 - ▶ Value function stores and reuses solutions.

Utility of DP in the RL context

DP is used for iterative **model-based** prediction and control in an MDP.

- ▶ Prediction:

- ▶ Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ and policy π
- ▶ Output: (estimated) value function $\hat{v}_\pi \approx v_\pi$

Utility of DP in the RL context

DP is used for iterative **model-based** prediction and control in an MDP.

- ▶ Prediction:

- ▶ Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ and policy π
- ▶ Output: (estimated) value function $\hat{v}_\pi \approx v_\pi$

- ▶ Control:

- ▶ Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
- ▶ Output: (estimated) optimal value function $\hat{v}_\pi^* \approx v_\pi^*$ or policy $\hat{\pi}^* \approx \pi^*$

Utility of DP in the RL context

DP is used for iterative **model-based** prediction and control in an MDP.

- ▶ Prediction:

- ▶ Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ and policy π
- ▶ Output: (estimated) value function $\hat{v}_\pi \approx v_\pi$

- ▶ Control:

- ▶ Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
- ▶ Output: (estimated) optimal value function $\hat{v}_\pi^* \approx v_\pi^*$ or policy $\hat{\pi}^* \approx \pi^*$

In both applications **DP requires full knowledge of the MDP** structure.

- ▶ Feasibility in real-world engineering applications (model vs. system) is therefore limited.

Utility of DP in the RL context

DP is used for iterative **model-based** prediction and control in an MDP.

- ▶ Prediction:

- ▶ Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ and policy π
- ▶ Output: (estimated) value function $\hat{v}_\pi \approx v_\pi$

- ▶ Control:

- ▶ Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
- ▶ Output: (estimated) optimal value function $\hat{v}_\pi^* \approx v_\pi^*$ or policy $\hat{\pi}^* \approx \pi^*$

In both applications **DP requires full knowledge of the MDP** structure.

- ▶ Feasibility in real-world engineering applications (model vs. system) is therefore limited.
- ▶ But: **following DP concepts are largely used in modern data-driven RL algorithms.**

Table of contents

- 1 Introduction
- 2 Policy evaluation (Prediction)**
- 3 Policy improvement (Control)
- 4 Policy and value iteration
- 5 Further aspects

Policy evaluation background (1)

- ▶ Problem: evaluate a given policy π to predict v_π .

Policy evaluation background (1)

- ▶ Problem: evaluate a given policy π to predict v_π .
- ▶ Recap: Bellman expectation equation for $s \in \mathcal{S}$ is given as

$$\begin{aligned} v_\pi(s) &= \mathbb{E}_\pi [G_t | S_t = s], \\ &= \mathbb{E}_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s], \\ &= \mathbb{E}_\pi [R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s]. \end{aligned}$$

Policy evaluation background (1)

- ▶ Problem: evaluate a given policy π to predict v_π .
- ▶ Recap: Bellman expectation equation for $s \in \mathcal{S}$ is given as

$$\begin{aligned}v_\pi(s) &= \mathbb{E}_\pi [G_t | S_t = s], \\&= \mathbb{E}_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s], \\&= \mathbb{E}_\pi [R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s].\end{aligned}$$

- ▶ Or in matrix form:

$$\begin{aligned}\mathbf{v}_\mathcal{S}^\pi &= \mathbf{r}_\mathcal{S}^\pi + \gamma \mathcal{P}_{ss'}^\pi \mathbf{v}_\mathcal{S}^\pi, \\ \begin{bmatrix} v_1^\pi \\ \vdots \\ v_n^\pi \end{bmatrix} &= \begin{bmatrix} \mathcal{R}_1^\pi \\ \vdots \\ \mathcal{R}_n^\pi \end{bmatrix} + \gamma \begin{bmatrix} p_{11}^\pi & \cdots & p_{1n}^\pi \\ \vdots & & \vdots \\ p_{n1}^\pi & \cdots & p_{nn}^\pi \end{bmatrix} \begin{bmatrix} v_1^\pi \\ \vdots \\ v_n^\pi \end{bmatrix}.\end{aligned}$$

Policy evaluation background (1)

- ▶ Problem: evaluate a given policy π to predict v_π .
- ▶ Recap: Bellman expectation equation for $s \in \mathcal{S}$ is given as

$$\begin{aligned}v_\pi(s) &= \mathbb{E}_\pi [G_t | S_t = s], \\&= \mathbb{E}_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s], \\&= \mathbb{E}_\pi [R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s].\end{aligned}$$

- ▶ Or in matrix form:

$$\begin{aligned}\mathbf{v}_\mathcal{S}^\pi &= \mathbf{r}_\mathcal{S}^\pi + \gamma \mathcal{P}_{ss'}^\pi \mathbf{v}_\mathcal{S}^\pi, \\ \begin{bmatrix} v_1^\pi \\ \vdots \\ v_n^\pi \end{bmatrix} &= \begin{bmatrix} \mathcal{R}_1^\pi \\ \vdots \\ \mathcal{R}_n^\pi \end{bmatrix} + \gamma \begin{bmatrix} p_{11}^\pi & \cdots & p_{1n}^\pi \\ \vdots & & \vdots \\ p_{n1}^\pi & \cdots & p_{nn}^\pi \end{bmatrix} \begin{bmatrix} v_1^\pi \\ \vdots \\ v_n^\pi \end{bmatrix}.\end{aligned}$$

- ▶ Solving the Bellman expectation equation for v_π requires handling a linear equation system with n unknowns (i.e., number of states).

Policy evaluation background (2)

- ▶ Problem: directly calculating v_π is numerically costly for high-dimensional state spaces (e.g., by matrix inversion).

Policy evaluation background (2)

- ▶ Problem: directly calculating v_π is numerically costly for high-dimensional state spaces (e.g., by matrix inversion).
- ▶ General idea: **apply iterative approximations** $\hat{v}_i(s) = v_i(s)$ of $v_\pi(s)$ with decreasing errors:

$$\|v_i(s) - v_\pi(s)\|_\infty \rightarrow 0 \quad \text{for } i = 1, 2, 3, \dots \quad (3.1)$$

Policy evaluation background (2)

- ▶ Problem: directly calculating v_π is numerically costly for high-dimensional state spaces (e.g., by matrix inversion).
- ▶ General idea: **apply iterative approximations** $\hat{v}_i(s) = v_i(s)$ of $v_\pi(s)$ with decreasing errors:

$$\|v_i(s) - v_\pi(s)\|_\infty \rightarrow 0 \quad \text{for } i = 1, 2, 3, \dots \quad (3.1)$$

- ▶ The Bellman equation in matrix form can be rewritten as:

$$\underbrace{(I - \gamma \mathcal{P}_{ss'}^\pi)}_A \underbrace{v_S^\pi}_x = \underbrace{r_S^\pi}_b. \quad (3.2)$$

Policy evaluation background (2)

- ▶ Problem: directly calculating v_π is numerically costly for high-dimensional state spaces (e.g., by matrix inversion).
- ▶ General idea: **apply iterative approximations** $\hat{v}_i(s) = v_i(s)$ of $v_\pi(s)$ with decreasing errors:

$$\|v_i(s) - v_\pi(s)\|_\infty \rightarrow 0 \quad \text{for } i = 1, 2, 3, \dots \quad (3.1)$$

- ▶ The Bellman equation in matrix form can be rewritten as:

$$\underbrace{(I - \gamma \mathcal{P}_{ss'}^\pi)}_A \underbrace{v_S^\pi}_x = \underbrace{r_S^\pi}_b. \quad (3.2)$$

- ▶ To iteratively solve this linear equation $Ax = b$, one can apply numerous methods such as
 - ▶ General gradient descent,
 - ▶ Richardson iteration,
 - ▶ Krylov subspace methods.

Richardson iteration (1)

In the MDP context, the Richardson iteration became the default solution approach to iteratively solve:

$$Ax = b.$$

The **Richardson iteration** is

$$x_{i+1} = x_i + \omega(b - Ax_i) \tag{3.3}$$

with ω being a scalar parameter that has to be chosen such that the sequence x_i converges.

Richardson iteration (1)

In the MDP context, the Richardson iteration became the default solution approach to iteratively solve:

$$\mathbf{A}\mathbf{x} = \mathbf{b}.$$

The **Richardson iteration** is

$$\mathbf{x}_{i+1} = \mathbf{x}_i + \omega(\mathbf{b} - \mathbf{A}\mathbf{x}_i) \quad (3.3)$$

with ω being a scalar parameter that has to be chosen such that the sequence \mathbf{x}_i converges. To choose ω we inspect the series of approximation errors $\mathbf{e}_i = \mathbf{x}_i - \mathbf{x}$ and apply it to (3.3):

$$\mathbf{e}_{i+1} = \mathbf{e}_i - \omega \mathbf{A}\mathbf{e}_i = (\mathbf{I} - \omega \mathbf{A}) \mathbf{e}_i. \quad (3.4)$$

Richardson iteration (1)

In the MDP context, the Richardson iteration became the default solution approach to iteratively solve:

$$\mathbf{A}\mathbf{x} = \mathbf{b}.$$

The **Richardson iteration** is

$$\mathbf{x}_{i+1} = \mathbf{x}_i + \omega(\mathbf{b} - \mathbf{A}\mathbf{x}_i) \quad (3.3)$$

with ω being a scalar parameter that has to be chosen such that the sequence \mathbf{x}_i converges. To choose ω we inspect the series of approximation errors $\mathbf{e}_i = \mathbf{x}_i - \mathbf{x}$ and apply it to (3.3):

$$\mathbf{e}_{i+1} = \mathbf{e}_i - \omega\mathbf{A}\mathbf{e}_i = (\mathbf{I} - \omega\mathbf{A})\mathbf{e}_i. \quad (3.4)$$

To evaluate convergence we inspect the following norm:

$$\|\mathbf{e}_{i+1}\|_{\infty} = \|(\mathbf{I} - \omega\mathbf{A})\mathbf{e}_i\|_{\infty}. \quad (3.5)$$

Richardson iteration (2)

Since any induced matrix norm is sub-multiplicative, we can approximate (3.5) by the inequality:

$$\|\mathbf{e}_{i+1}\|_{\infty} \leq \|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} \|\mathbf{e}_i\|_{\infty}. \quad (3.6)$$

Richardson iteration (2)

Since any induced matrix norm is sub-multiplicative, we can approximate (3.5) by the inequality:

$$\|e_{i+1}\|_{\infty} \leq \|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} \|e_i\|_{\infty}. \quad (3.6)$$

Hence, the series converges if

$$\|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} < 1. \quad (3.7)$$

Richardson iteration (2)

Since any induced matrix norm is sub-multiplicative, we can approximate (3.5) by the inequality:

$$\|e_{i+1}\|_{\infty} \leq \|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} \|e_i\|_{\infty}. \quad (3.6)$$

Hence, the series converges if

$$\|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} < 1. \quad (3.7)$$

Inserting from (3.2) leads to:

$$\|(\mathbf{I}(1 - \omega) + \omega\gamma\mathcal{P}_{ss'}^{\pi})\|_{\infty} < 1. \quad (3.8)$$

Richardson iteration (2)

Since any induced matrix norm is sub-multiplicative, we can approximate (3.5) by the inequality:

$$\|e_{i+1}\|_{\infty} \leq \|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} \|e_i\|_{\infty}. \quad (3.6)$$

Hence, the series converges if

$$\|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} < 1. \quad (3.7)$$

Inserting from (3.2) leads to:

$$\|(\mathbf{I}(1 - \omega) + \omega \gamma \mathcal{P}_{ss'}^{\pi})\|_{\infty} < 1. \quad (3.8)$$

For $\omega = 1$ we receive:

$$\gamma \|(\mathcal{P}_{ss'}^{\pi})\|_{\infty} < 1. \quad (3.9)$$

Richardson iteration (2)

Since any induced matrix norm is sub-multiplicative, we can approximate (3.5) by the inequality:

$$\|e_{i+1}\|_{\infty} \leq \|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} \|e_i\|_{\infty}. \quad (3.6)$$

Hence, the series converges if

$$\|(\mathbf{I} - \omega \mathbf{A})\|_{\infty} < 1. \quad (3.7)$$

Inserting from (3.2) leads to:

$$\|(\mathbf{I}(1 - \omega) + \omega\gamma\mathcal{P}_{ss'}^{\pi})\|_{\infty} < 1. \quad (3.8)$$

For $\omega = 1$ we receive:

$$\gamma \|(\mathcal{P}_{ss'}^{\pi})\|_{\infty} < 1. \quad (3.9)$$

Since the row elements of $\mathcal{P}_{ss'}^{\pi}$ always sum up to 1,

$$\gamma < 1 \quad (3.10)$$

follows. Hence, **when discounting the Richardson iteration always converges for MDPs** even if we assume $\omega = 1$.

Iterative policy evaluation by Richardson iteration (1)

Applying the Richardson iteration (3.3) with $w = 1$ to the Bellman equation for any $s \in \mathcal{S}$ at iteration i results in:

$$v_{i+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_i(s') \right). \quad (3.11)$$

Iterative policy evaluation by Richardson iteration (1)

Applying the Richardson iteration (3.3) with $w = 1$ to the Bellman equation for any $s \in \mathcal{S}$ at iteration i results in:

$$v_{i+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_i(s') \right). \quad (3.11)$$

Matrix form then is:

$$\mathbf{v}_{\mathcal{S},i+1}^{\pi} = \mathbf{r}_{\mathcal{S}}^{\pi} + \gamma \mathcal{P}_{ss'}^{\pi} \mathbf{v}_{\mathcal{S},i}^{\pi}. \quad (3.12)$$

Iterative policy evaluation by Richardson iteration (1)

Applying the Richardson iteration (3.3) with $w = 1$ to the Bellman equation for any $s \in \mathcal{S}$ at iteration i results in:

$$v_{i+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_i(s') \right). \quad (3.11)$$

Matrix form then is:

$$\mathbf{v}_{\mathcal{S},i+1}^{\pi} = \mathbf{r}_{\mathcal{S}}^{\pi} + \gamma \mathcal{P}_{ss'}^{\pi} \mathbf{v}_{\mathcal{S},i}^{\pi}. \quad (3.12)$$




Fig. 3.1: Backup diagram for iterative policy evaluation

Iterative policy evaluation by Richardson iteration (2)

- ▶ During one Richardson iteration the 'old' value of s is replaced with a 'new' value from the 'old' values of the successor state s' .
 - ▶ Update $v_{i+1}(s)$ from $v_i(s')$, see Fig. 3.1.
 - ▶ Updating estimates (v_{i+1}) on the basis of other estimates (v_i) is often called **bootstrapping**.

Iterative policy evaluation by Richardson iteration (2)

- ▶ During one Richardson iteration the 'old' value of s is replaced with a 'new' value from the 'old' values of the successor state s' .
 - ▶ Update $v_{i+1}(s)$ from $v_i(s')$, see Fig. 3.1.
 - ▶ Updating estimates (v_{i+1}) on the basis of other estimates (v_i) is often called **bootstrapping**.
- ▶ The Richardson iteration can be interpreted as a gradient descent algorithm for solving (3.2).

Iterative policy evaluation by Richardson iteration (2)

- ▶ During one Richardson iteration the 'old' value of s is replaced with a 'new' value from the 'old' values of the successor state s' .
 - ▶ Update $v_{i+1}(s)$ from $v_i(s')$, see Fig. 3.1.
 - ▶ Updating estimates (v_{i+1}) on the basis of other estimates (v_i) is often called **bootstrapping**.
- ▶ The Richardson iteration can be interpreted as a gradient descent algorithm for solving (3.2).
- ▶ This leads to **synchronous, full backups** of the entire state space \mathcal{S} .

Iterative policy evaluation by Richardson iteration (2)

- ▶ During one Richardson iteration the 'old' value of s is replaced with a 'new' value from the 'old' values of the successor state s' .
 - ▶ Update $v_{i+1}(s)$ from $v_i(s')$, see Fig. 3.1.
 - ▶ Updating estimates (v_{i+1}) on the basis of other estimates (v_i) is often called **bootstrapping**.
- ▶ The Richardson iteration can be interpreted as a gradient descent algorithm for solving (3.2).
- ▶ This leads to **synchronous, full backups** of the entire state space \mathcal{S} .
- ▶ Also called **expected update** because it is based on the expectation over all possible next states (utilizing full model knowledge).

Iterative policy evaluation by Richardson iteration (2)

- ▶ During one Richardson iteration the 'old' value of s is replaced with a 'new' value from the 'old' values of the successor state s' .
 - ▶ Update $v_{i+1}(s)$ from $v_i(s')$, see Fig. 3.1.
 - ▶ Updating estimates (v_{i+1}) on the basis of other estimates (v_i) is often called **bootstrapping**.
- ▶ The Richardson iteration can be interpreted as a gradient descent algorithm for solving (3.2).
- ▶ This leads to **synchronous, full backups** of the entire state space \mathcal{S} .
- ▶ Also called **expected update** because it is based on the expectation over all possible next states (utilizing full model knowledge).
- ▶ In subsequent lectures, the expected update will be supplemented by data-driven samples from the environment.

Iterative policy evaluation example: forest tree MDP

Let's reuse the forest tree MDP example with *fifty-fifty policy*:

$$\mathcal{P}_{ss'}^{\pi} = \begin{bmatrix} 0 & \frac{1-\alpha}{2} & 0 & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & 0 & 1 \end{bmatrix}, \quad \mathbf{r}_S^{\pi} = \begin{bmatrix} 0.5 \\ 1 \\ 2 \\ 0 \end{bmatrix}.$$

Iterative policy evaluation example: forest tree MDP

Let's reuse the forest tree MDP example with *fifty-fifty policy*:

$$\mathcal{P}_{ss'}^{\pi} = \begin{bmatrix} 0 & \frac{1-\alpha}{2} & 0 & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & 0 & 1 \end{bmatrix}, \quad \mathbf{r}_{\mathcal{S}}^{\pi} = \begin{bmatrix} 0.5 \\ 1 \\ 2 \\ 0 \end{bmatrix}.$$

i	$v_i(s=1)$	$v_i(s=2)$	$v_i(s=3)$	$v_i(s=4)$
0	0	0	0	0

Iterative policy evaluation example: forest tree MDP

Let's reuse the forest tree MDP example with *fifty-fifty policy*:

$$\mathcal{P}_{ss'}^{\pi} = \begin{bmatrix} 0 & \frac{1-\alpha}{2} & 0 & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & 0 & 1 \end{bmatrix}, \quad \mathbf{r}_{\mathcal{S}}^{\pi} = \begin{bmatrix} 0.5 \\ 1 \\ 2 \\ 0 \end{bmatrix}.$$

i	$v_i(s=1)$	$v_i(s=2)$	$v_i(s=3)$	$v_i(s=4)$
0	0	0	0	0
1	0.5	1	2	0

Iterative policy evaluation example: forest tree MDP

Let's reuse the forest tree MDP example with *fifty-fifty policy*:

$$\mathcal{P}_{ss'}^{\pi} = \begin{bmatrix} 0 & \frac{1-\alpha}{2} & 0 & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & 0 & 1 \end{bmatrix}, \quad \mathbf{r}_{\mathcal{S}}^{\pi} = \begin{bmatrix} 0.5 \\ 1 \\ 2 \\ 0 \end{bmatrix}.$$

i	$v_i(s=1)$	$v_i(s=2)$	$v_i(s=3)$	$v_i(s=4)$
0	0	0	0	0
1	0.5	1	2	0
2	0.82	1.64	2.64	0

Iterative policy evaluation example: forest tree MDP

Let's reuse the forest tree MDP example with *fifty-fifty policy*:

$$\mathcal{P}_{ss'}^{\pi} = \begin{bmatrix} 0 & \frac{1-\alpha}{2} & 0 & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & 0 & 1 \end{bmatrix}, \quad \mathbf{r}_{\mathcal{S}}^{\pi} = \begin{bmatrix} 0.5 \\ 1 \\ 2 \\ 0 \end{bmatrix}.$$

i	$v_i(s=1)$	$v_i(s=2)$	$v_i(s=3)$	$v_i(s=4)$
0	0	0	0	0
1	0.5	1	2	0
2	0.82	1.64	2.64	0
3	1.03	1.85	2.85	0

Iterative policy evaluation example: forest tree MDP

Let's reuse the forest tree MDP example with *fifty-fifty policy*:

$$\mathcal{P}_{ss'}^{\pi} = \begin{bmatrix} 0 & \frac{1-\alpha}{2} & 0 & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & \frac{1-\alpha}{2} & \frac{1+\alpha}{2} \\ 0 & 0 & 0 & 1 \end{bmatrix}, \quad \mathbf{r}_S^{\pi} = \begin{bmatrix} 0.5 \\ 1 \\ 2 \\ 0 \end{bmatrix}.$$

i	$v_i(s=1)$	$v_i(s=2)$	$v_i(s=3)$	$v_i(s=4)$
0	0	0	0	0
1	0.5	1	2	0
2	0.82	1.64	2.64	0
3	1.03	1.85	2.85	0
\vdots	\vdots	\vdots	\vdots	\vdots
∞	1.12	1.94	2.94	0

Tab. 3.1: Policy evaluation by Richardson iteration (3.12) for forest tree MDP with $\gamma = 0.8$ and $\alpha = 0.2$

Variant: in-place updates

Instead of applying (3.12) to the entire vector $v_{\mathcal{S},i+1}^{\pi}$ in 'one shot' (synchronous backup), an elementwise **in-place** version of the policy evaluation can be carried out:

input: full model of the MDP, i.e., $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ including policy π

parameter: $\delta > 0$ as accuracy termination threshold

init: $v_0(s) \forall s \in \mathcal{S}$ arbitrary except $v_0(s) = 0$ if s is terminal

repeat

$\Delta \leftarrow 0;$

for $\forall s \in \mathcal{S}$ **do**

$\tilde{v} \leftarrow \hat{v}(s);$

$\hat{v}(s) \leftarrow \sum_{a \in \mathcal{A}} \pi(a|s) (\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a \hat{v}(s'));$

$\Delta \leftarrow \max(\Delta, |\tilde{v} - \hat{v}(s)|);$

until $\Delta < \delta;$

Algo. 3.1: Iterative policy evaluation using in-place updates (output: estimate of $v_{\mathcal{S}}^{\pi}$)

In-place policy evaluation updates for forest tree MDP

- ▶ In-place algorithms allow to update states in a beneficial order.

In-place policy evaluation updates for forest tree MDP

- ▶ In-place algorithms allow to update states in a beneficial order.
- ▶ May converge faster than regular Richardson iteration if state update order is chosen wisely (sweep through state space).

In-place policy evaluation updates for forest tree MDP

- ▶ In-place algorithms allow to update states in a beneficial order.
- ▶ May converge faster than regular Richardson iteration if state update order is chosen wisely (sweep through state space).
- ▶ For forest tree MDP: reverse order, i.e., start with $s = 4$.

In-place policy evaluation updates for forest tree MDP

- ▶ In-place algorithms allow to update states in a beneficial order.
- ▶ May converge faster than regular Richardson iteration if state update order is chosen wisely (sweep through state space).
- ▶ For forest tree MDP: reverse order, i.e., start with $s = 4$.
- ▶ As can be seen in Tab. 3.2 the in-place updates especially converge faster for the 'early states'.

i	$v_i(s = 1)$	$v_i(s = 2)$	$v_i(s = 3)$	$v_i(s = 4)$
0	0	0	0	0
1	1.03	1.64	2	0
2	1.09	1.85	2.64	0
3	1.11	1.91	2.85	0
\vdots	\vdots	\vdots	\vdots	\vdots
∞	1.12	1.94	2.94	0

Tab. 3.2: In-place updates for forest tree MDP

Table of contents

- 1 Introduction
- 2 Policy evaluation (Prediction)
- 3 Policy improvement (Control)**
- 4 Policy and value iteration
- 5 Further aspects

General idea on policy improvement

- ▶ If we know v_π of a given MDP, how to improve the policy?

General idea on policy improvement

- ▶ If we know v_π of a given MDP, how to improve the policy?
- ▶ The simple idea of policy improvement is:
 - ▶ Consider a new (non-policy conform) action $a \neq \pi(s)$.
 - ▶ Follow thereafter the current policy π .
 - ▶ Check the action value of this 'new move'. If it is better than the 'old' value, take it:

$$q_\pi(s, a) = \mathbb{E} [R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s, A_t = a] . \quad (3.13)$$

General idea on policy improvement

- ▶ If we know v_π of a given MDP, how to improve the policy?
- ▶ The simple idea of policy improvement is:
 - ▶ Consider a new (non-policy conform) action $a \neq \pi(s)$.
 - ▶ Follow thereafter the current policy π .
 - ▶ Check the action value of this 'new move'. If it is better than the 'old' value, take it:

$$q_\pi(s, a) = \mathbb{E} [R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s, A_t = a] . \quad (3.13)$$

Theorem 3.1: Policy improvement

If for any deterministic policy pair π and π'

$$q_\pi(s, \pi'(s)) \geq v_\pi(s) \quad \forall s \in \mathcal{S} \quad (3.14)$$

applies, then the policy π' must be as good as or better than π . Hence, it obtains greater or equal expected return

$$v_{\pi'}(s) \geq v_\pi(s) \quad \forall s \in \mathcal{S}. \quad (3.15)$$

Proof of policy improvement theorem

Start with (3.14) and recursively reapply (3.13):

$$v_{\pi}(s) \leq q_{\pi}(s, \pi'(s)),$$

(3.16)

Proof of policy improvement theorem

Start with (3.14) and recursively reapply (3.13):

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)), \\ &= \mathbb{E} \left[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = \pi'(s) \right], \end{aligned}$$

(3.16)

Proof of policy improvement theorem

Start with (3.14) and recursively reapply (3.13):

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)), \\ &= \mathbb{E} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = \pi'(s)] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] , \end{aligned} \tag{3.16}$$

Proof of policy improvement theorem

Start with (3.14) and recursively reapply (3.13):

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)), \\ &= \mathbb{E} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = \pi'(s)] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] , \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) | S_t = s] , \end{aligned} \tag{3.16}$$

Proof of policy improvement theorem

Start with (3.14) and recursively reapply (3.13):

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)), \\ &= \mathbb{E} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = \pi'(s)] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] , \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) | S_t = s] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma \mathbb{E}_{\pi'} [R_{t+2} + \gamma v_{\pi}(S_{t+2}) | S_{t+1}, \pi'(S_{t+1})] | S_t = s] , \end{aligned} \tag{3.16}$$

Proof of policy improvement theorem

Start with (3.14) and recursively reapply (3.13):

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)), \\ &= \mathbb{E} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = \pi'(s)] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] , \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) | S_t = s] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma \mathbb{E}_{\pi'} [R_{t+2} + \gamma v_{\pi}(S_{t+2}) | S_{t+1}, \pi'(S_{t+1})] | S_t = s] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 v_{\pi}(S_{t+2}) | S_t = s] , \end{aligned} \tag{3.16}$$

Proof of policy improvement theorem

Start with (3.14) and recursively reapply (3.13):

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)), \\ &= \mathbb{E} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = \pi'(s)] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] , \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) | S_t = s] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma \mathbb{E}_{\pi'} [R_{t+2} + \gamma v_{\pi}(S_{t+2}) | S_{t+1}, \pi'(S_{t+1})] | S_t = s] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 v_{\pi}(S_{t+2}) | S_t = s] , \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 v_{\pi}(S_{t+3}) | S_t = s] , \end{aligned} \tag{3.16}$$

Proof of policy improvement theorem

Start with (3.14) and recursively reapply (3.13):

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)), \\ &= \mathbb{E} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = \pi'(s)] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] , \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) | S_t = s] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma \mathbb{E}_{\pi'} [R_{t+2} + \gamma v_{\pi}(S_{t+2}) | S_{t+1}, \pi'(S_{t+1})] | S_t = s] , \\ &= \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 v_{\pi}(S_{t+2}) | S_t = s] , \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 v_{\pi}(S_{t+3}) | S_t = s] , \\ &\vdots \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \cdots | S_t = s] , \\ &= v_{\pi'}(s). \end{aligned} \tag{3.16}$$

Greedy policy improvement (1)

- ▶ So far, policy improvement addressed only changing the policy at a single state.

Greedy policy improvement (1)

- ▶ So far, policy improvement addressed only changing the policy at a single state.
- ▶ Now, extend this scheme to all states by selecting the best action according to $q_\pi(s, a)$ in every state (**greedy policy improvement**):

Greedy policy improvement (1)

- ▶ So far, policy improvement addressed only changing the policy at a single state.
- ▶ Now, extend this scheme to all states by selecting the best action according to $q_\pi(s, a)$ in every state (**greedy policy improvement**):

$$\begin{aligned}\pi'(s) &= \arg \max_{a \in \mathcal{A}} q_\pi(s, a), \\ &= \arg \max_{a \in \mathcal{A}} \mathbb{E} [R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s, A_t = a], \\ &= \arg \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_\pi(s').\end{aligned}\tag{3.17}$$

Greedy policy improvement (2)

- ▶ Each greedy policy improvement takes the best action in a one-step look-ahead search and, therefore, satisfies Theo. 3.1.

Greedy policy improvement (2)

- ▶ Each greedy policy improvement takes the best action in a one-step look-ahead search and, therefore, satisfies Theo. 3.1.
- ▶ If after a policy improvement step $v_\pi(s) = v_{\pi'}(s)$ applies, it follows:

$$\begin{aligned} v_{\pi'}(s) &= \max_{a \in \mathcal{A}} \mathbb{E} [R_{t+1} + \gamma v_{\pi'}(S_{t+1}) | S_t = s, A_t = a], \\ &= \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_{\pi'}(s'). \end{aligned} \tag{3.18}$$

Greedy policy improvement (2)

- ▶ Each greedy policy improvement takes the best action in a one-step look-ahead search and, therefore, satisfies Theo. 3.1.
- ▶ If after a policy improvement step $v_\pi(s) = v_{\pi'}(s)$ applies, it follows:

$$\begin{aligned} v_{\pi'}(s) &= \max_{a \in \mathcal{A}} \mathbb{E} [R_{t+1} + \gamma v_{\pi'}(S_{t+1}) | S_t = s, A_t = a], \\ &= \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_{\pi'}(s'). \end{aligned} \tag{3.18}$$

- ▶ This is the Bellman optimality equation, which guarantees that $\pi' = \pi$ must be optimal policies.

Greedy policy improvement (2)

- ▶ Each greedy policy improvement takes the best action in a one-step look-ahead search and, therefore, satisfies Theo. 3.1.
- ▶ If after a policy improvement step $v_{\pi}(s) = v_{\pi'}(s)$ applies, it follows:

$$\begin{aligned} v_{\pi'}(s) &= \max_{a \in \mathcal{A}} \mathbb{E} [R_{t+1} + \gamma v_{\pi'}(S_{t+1}) | S_t = s, A_t = a], \\ &= \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_{\pi'}(s'). \end{aligned} \tag{3.18}$$

- ▶ This is the Bellman optimality equation, which guarantees that $\pi' = \pi$ must be optimal policies.
- ▶ Although the proof for policy improvement theorem was presented for deterministic policies, transfer to stochastic policies $\pi(a|s)$ is possible.

Greedy policy improvement (2)

- ▶ Each greedy policy improvement takes the best action in a one-step look-ahead search and, therefore, satisfies Theo. 3.1.
- ▶ If after a policy improvement step $v_{\pi}(s) = v_{\pi'}(s)$ applies, it follows:

$$\begin{aligned} v_{\pi'}(s) &= \max_{a \in \mathcal{A}} \mathbb{E} [R_{t+1} + \gamma v_{\pi'}(S_{t+1}) | S_t = s, A_t = a], \\ &= \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_{\pi'}(s'). \end{aligned} \tag{3.18}$$

- ▶ This is the Bellman optimality equation, which guarantees that $\pi' = \pi$ must be optimal policies.
- ▶ Although the proof for policy improvement theorem was presented for deterministic policies, transfer to stochastic policies $\pi(a|s)$ is possible.
- ▶ Takeaway message: **policy improvement theorem guarantees finding optimal policies in finite MDPs** (e.g., by DP).

Table of contents

- 1 Introduction
- 2 Policy evaluation (Prediction)
- 3 Policy improvement (Control)
- 4 Policy and value iteration**
- 5 Further aspects

Concept of policy iteration

- Policy iteration **combines the previous policy evaluation and policy improvement** in an iterative sequence:

$$\pi_0 \rightarrow v_{\pi_0} \rightarrow \pi_1 \rightarrow v_{\pi_1} \rightarrow \cdots \pi^* \rightarrow v_{\pi^*} \quad (3.19)$$

- Evaluate \rightarrow improve \rightarrow evaluate \rightarrow improve ...


Concept of policy iteration

- ▶ Policy iteration **combines the previous policy evaluation and policy improvement** in an iterative sequence:

$$\pi_0 \rightarrow v_{\pi_0} \rightarrow \pi_1 \rightarrow v_{\pi_1} \rightarrow \cdots \pi^* \rightarrow v_{\pi^*} \quad (3.19)$$

- ▶ Evaluate \rightarrow improve \rightarrow evaluate \rightarrow improve ...
- ▶ In the 'classic' policy iteration, each policy evaluation step in (3.19) is fully executed, i.e., for each policy π_i an exact estimate of v_{π_i} is provided either by iterative policy evaluation with a sufficiently high number of steps or by any other method that fully solves (3.2).

Policy iteration example: forest tree MDP (1)



- ▶ Two actions possible in each state:
 - ▶ Wait $a = w$: let the tree grow.
 - ▶ Cut $a = c$: gather the wood.

Policy iteration example: forest tree MDP (2)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree hater' initial policy:

Policy iteration example: forest tree MDP (2)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree hater' initial policy:

① $\pi_0 = \pi(a = c|s) \quad \forall s \in \mathcal{S}.$

Policy iteration example: forest tree MDP (2)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree hater' initial policy:

- 1 $\pi_0 = \pi(a = c|s) \quad \forall s \in \mathcal{S}.$
- 2 Policy evaluation: $v_{\mathcal{S}}^{\pi_0} = [1 \quad 2 \quad 3 \quad 0]^T$

Policy iteration example: forest tree MDP (2)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree hater' initial policy:

- ❶ $\pi_0 = \pi(a = c|s) \quad \forall s \in \mathcal{S}.$
- ❷ Policy evaluation: $v_{\mathcal{S}}^{\pi_0} = [1 \quad 2 \quad 3 \quad 0]^T$
- ❸ Greedy policy improvement:

$$\begin{aligned}\pi_1(s) &= \arg \max_{a \in \mathcal{A}} \mathbb{E}[R_{t+1} + \gamma v_{\pi_0}(S_{t+1}) | S_t = s, A_t = a], \\ &= \{\pi(a = w | s = 1), \pi(a = c | s = 2), \pi(a = c | s = 3)\}\end{aligned}$$

Policy iteration example: forest tree MDP (2)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree hater' initial policy:

- ① $\pi_0 = \pi(a = c|s) \quad \forall s \in \mathcal{S}.$
- ② Policy evaluation: $v_{\mathcal{S}}^{\pi_0} = [1 \quad 2 \quad 3 \quad 0]^T$
- ③ Greedy policy improvement:

$$\begin{aligned}\pi_1(s) &= \arg \max_{a \in \mathcal{A}} \mathbb{E}[R_{t+1} + \gamma v_{\pi_0}(S_{t+1}) | S_t = s, A_t = a], \\ &= \{\pi(a = w|s = 1), \pi(a = c|s = 2), \pi(a = c|s = 3)\}\end{aligned}$$

- ④ Policy evaluation: $v_{\mathcal{S}}^{\pi_1} = [1.28 \quad 2 \quad 3 \quad 0]^T$

Policy iteration example: forest tree MDP (2)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree hater' initial policy:

- 1 $\pi_0 = \pi(a = c|s) \quad \forall s \in \mathcal{S}.$
- 2 Policy evaluation: $v_S^{\pi_0} = [1 \quad 2 \quad 3 \quad 0]^T$
- 3 Greedy policy improvement:

$$\begin{aligned}\pi_1(s) &= \arg \max_{a \in \mathcal{A}} \mathbb{E}[R_{t+1} + \gamma v_{\pi_0}(S_{t+1}) | S_t = s, A_t = a], \\ &= \{\pi(a = w | s = 1), \pi(a = c | s = 2), \pi(a = c | s = 3)\}\end{aligned}$$

- 4 Policy evaluation: $v_S^{\pi_1} = [1.28 \quad 2 \quad 3 \quad 0]^T$
- 5 Greedy policy improvement:

$$\begin{aligned}\pi_2(s) &= \arg \max_{a \in \mathcal{A}} \mathbb{E}[R_{t+1} + \gamma v_{\pi_1}(S_{t+1}) | S_t = s, A_t = a], \\ &= \{\pi(a = w | s = 1), \pi(a = c | s = 2), \pi(a = c | s = 3)\}, \\ &= \pi_1(s) \\ &= \pi^*\end{aligned}$$

Policy iteration example: forest tree MDP (3)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree lover' initial policy:

Policy iteration example: forest tree MDP (3)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree lover' initial policy:

① $\pi_0 = \pi(a = w|s) \quad \forall s \in \mathcal{S}.$

Policy iteration example: forest tree MDP (3)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree lover' initial policy:

- ① $\pi_0 = \pi(a = w|s) \quad \forall s \in \mathcal{S}.$
- ② Policy evaluation: $v_{\mathcal{S}}^{\pi_0} = [1.14 \quad 1.78 \quad 2.78 \quad 0]^T$

Policy iteration example: forest tree MDP (3)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree lover' initial policy:

- ① $\pi_0 = \pi(a = w|s) \quad \forall s \in \mathcal{S}.$
- ② Policy evaluation: $v_{\mathcal{S}}^{\pi_0} = [1.14 \quad 1.78 \quad 2.78 \quad 0]^T$
- ③ Greedy policy improvement:

$$\begin{aligned}\pi_1(s) &= \arg \max_{a \in \mathcal{A}} \mathbb{E}[R_{t+1} + \gamma v_{\pi_0}(S_{t+1}) | S_t = s, A_t = a], \\ &= \{\pi(a = w|s = 1), \pi(a = c|s = 2), \pi(a = c|s = 3)\}\end{aligned}$$

Policy iteration example: forest tree MDP (3)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree lover' initial policy:

① $\pi_0 = \pi(a = w|s) \quad \forall s \in \mathcal{S}.$

② Policy evaluation: $v_S^{\pi_0} = [1.14 \quad 1.78 \quad 2.78 \quad 0]^T$

③ Greedy policy improvement:

$$\begin{aligned}\pi_1(s) &= \arg \max_{a \in \mathcal{A}} \mathbb{E}[R_{t+1} + \gamma v_{\pi_0}(S_{t+1}) | S_t = s, A_t = a], \\ &= \{\pi(a = w|s = 1), \pi(a = c|s = 2), \pi(a = c|s = 3)\}\end{aligned}$$

④ Policy evaluation: $v_S^{\pi_1} = [1.28 \quad 2 \quad 3 \quad 0]^T$

Policy iteration example: forest tree MDP (3)

Assume $\alpha = 0.2$ and $\gamma = 0.8$ and start with 'tree lover' initial policy:

① $\pi_0 = \pi(a = w|s) \quad \forall s \in \mathcal{S}.$

② Policy evaluation: $v_S^{\pi_0} = [1.14 \quad 1.78 \quad 2.78 \quad 0]^T$

③ Greedy policy improvement:

$$\begin{aligned}\pi_1(s) &= \arg \max_{a \in \mathcal{A}} \mathbb{E}[R_{t+1} + \gamma v_{\pi_0}(S_{t+1}) | S_t = s, A_t = a], \\ &= \{\pi(a = w|s = 1), \pi(a = c|s = 2), \pi(a = c|s = 3)\}\end{aligned}$$

④ Policy evaluation: $v_S^{\pi_1} = [1.28 \quad 2 \quad 3 \quad 0]^T$

⑤ Greedy policy improvement:

$$\begin{aligned}\pi_2(s) &= \arg \max_{a \in \mathcal{A}} \mathbb{E}[R_{t+1} + \gamma v_{\pi_1}(S_{t+1}) | S_t = s, A_t = a], \\ &= \{\pi(a = w|s = 1), \pi(a = c|s = 2), \pi(a = c|s = 3)\}, \\ &= \pi_1(s) \\ &= \pi^*\end{aligned}$$

Value iteration (1)

- ▶ Policy iteration involves full policy evaluation steps between policy improvements.
- ▶ In large state-space MDPs the full policy evaluation may be numerically very costly.

Value iteration (1)

- ▶ Policy iteration involves full policy evaluation steps between policy improvements.
- ▶ In large state-space MDPs the full policy evaluation may be numerically very costly.
- ▶ **Value iteration**: One step iterative policy evaluation followed by policy improvement.

Value iteration (1)

- ▶ Policy iteration involves full policy evaluation steps between policy improvements.
- ▶ In large state-space MDPs the full policy evaluation may be numerically very costly.
- ▶ **Value iteration**: One step iterative policy evaluation followed by policy improvement.
- ▶ Allows simple update rule which **combines policy improvement with truncated policy evaluation in a single step**:

$$\begin{aligned} v_{i+1}(s) &= \max_{a \in \mathcal{A}} \mathbb{E} [R_{t+1} + \gamma v_i(S_{t+1}) | S_t = s, A_t = a], \\ &= \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a v_i(s'). \end{aligned} \tag{3.20}$$

Value iteration (2)

```
input: full model of the MDP, i.e.,  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$   
parameter:  $\delta > 0$  as accuracy termination threshold  
init:  $v_0(s) \forall s \in \mathcal{S}$  arbitrary except  $v_0(s) = 0$  if  $s$  is terminal  
repeat  
     $\Delta \leftarrow 0$ ;  
    for  $\forall s \in \mathcal{S}$  do  
         $\tilde{v} \leftarrow \hat{v}(s)$ ;  
         $\hat{v}(s) \leftarrow \max_{a \in \mathcal{A}} (\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a \hat{v}(s'))$ ;  
         $\Delta \leftarrow \max(\Delta, |\tilde{v} - \hat{v}(s)|)$ ;  
until  $\Delta < \delta$ ;  
output: deterministic policy  $\pi \approx \pi^*$ , such that  
 $\pi(s) \leftarrow \arg \max_{a \in \mathcal{A}} (\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}^a \hat{v}(s'))$ ;
```

Algo. 3.2: Value iteration (note: compared to policy iteration, value iteration does not require an initial policy but only a state-value guess)

Value iteration example: forest tree MDP

- ▶ Assume again $\alpha = 0.2$ and $\gamma = 0.8$.

Value iteration example: forest tree MDP

- ▶ Assume again $\alpha = 0.2$ and $\gamma = 0.8$.
- ▶ Similar to in-place update policy evaluation, reverse order and start value iteration with $s = 4$.

Value iteration example: forest tree MDP

- ▶ Assume again $\alpha = 0.2$ and $\gamma = 0.8$.
- ▶ Similar to in-place update policy evaluation, reverse order and start value iteration with $s = 4$.
- ▶ As shown in Tab. 3.3 value iteration converges in one step (for the given problem) to the optimal state value.

i	$v_i(s = 1)$	$v_i(s = 2)$	$v_i(s = 3)$	$v_i(s = 4)$
0	0	0	0	0
1	1.28	2	3	0
*	1.28	2	3	0

Tab. 3.3: Value iteration for forest tree MDP

Table of contents

- 1 Introduction
- 2 Policy evaluation (Prediction)
- 3 Policy improvement (Control)
- 4 Policy and value iteration
- 5 Further aspects

Summarizing DP algorithms

- ▶ All DP algorithms are based on the state value $v(s)$.
 - ▶ Complexity is $\mathcal{O}(m \cdot n^2)$ for m actions and n states.
 - ▶ Evaluate all n^2 state transitions while considering up to m actions per state.

Summarizing DP algorithms

- ▶ All DP algorithms are based on the state value $v(s)$.
 - ▶ Complexity is $\mathcal{O}(m \cdot n^2)$ for m actions and n states.
 - ▶ Evaluate all n^2 state transitions while considering up to m actions per state.
- ▶ Could be also applied to action values $q(s, a)$.
 - ▶ Complexity is inferior with $\mathcal{O}(m^2 \cdot n^2)$.
 - ▶ There are up to m^2 action values which require n^2 state transition evaluations each.

Summarizing DP algorithms

- ▶ All DP algorithms are based on the state value $v(s)$.
 - ▶ Complexity is $\mathcal{O}(m \cdot n^2)$ for m actions and n states.
 - ▶ Evaluate all n^2 state transitions while considering up to m actions per state.
- ▶ Could be also applied to action values $q(s, a)$.
 - ▶ Complexity is inferior with $\mathcal{O}(m^2 \cdot n^2)$.
 - ▶ There are up to m^2 action values which require n^2 state transition evaluations each.

Problem	Relevant Equations	Algorithm
prediction	Bellman expectation eq.	policy evaluation
control	Bellman expectation eq. & greedy policy improvement	policy iteration
control	Bellman optimality eq.	value iteration

Tab. 3.4: Short overview addressing the treated DP algorithms

Curse of dimensionality

- ▶ DP is much more efficient than an exhaustive search over all n states and m actions in finite MDPs in order to find an optimal policy.
 - ▶ Exhaustive search for deterministic policies: m^n evaluations.
 - ▶ DP results in polynomial complexity regarding m and n .

Curse of dimensionality

- ▶ DP is much more efficient than an exhaustive search over all n states and m actions in finite MDPs in order to find an optimal policy.
 - ▶ Exhaustive search for deterministic policies: m^n evaluations.
 - ▶ DP results in polynomial complexity regarding m and n .
- ▶ Nevertheless, DP uses full-width backups:
 - ▶ For each state update, every successor state and action is considered.
 - ▶ While utilizing full knowledge of the MDP structure.

Curse of dimensionality

- ▶ DP is much more efficient than an exhaustive search over all n states and m actions in finite MDPs in order to find an optimal policy.
 - ▶ Exhaustive search for deterministic policies: m^n evaluations.
 - ▶ DP results in polynomial complexity regarding m and n .
- ▶ Nevertheless, DP uses full-width backups:
 - ▶ For each state update, every successor state and action is considered.
 - ▶ While utilizing full knowledge of the MDP structure.
- ▶ Hence, DP is can be effective up to medium-sized MDPs (i.e., million finite states)

Curse of dimensionality

- ▶ DP is much more efficient than an exhaustive search over all n states and m actions in finite MDPs in order to find an optimal policy.
 - ▶ Exhaustive search for deterministic policies: m^n evaluations.
 - ▶ DP results in polynomial complexity regarding m and n .
- ▶ Nevertheless, DP uses full-width backups:
 - ▶ For each state update, every successor state and action is considered.
 - ▶ While utilizing full knowledge of the MDP structure.
- ▶ Hence, DP is can be effective up to medium-sized MDPs (i.e., million finite states)
- ▶ For large problems DP suffers from the **curse of dimensionality**:
 - ▶ Single update step may become computational infeasible.
 - ▶ Also: if continuous states need quantization, number of finite states n grows exponentially with the number of state variables (assuming fixed number of discretization levels).

Generalized policy iteration (GPI)

- ▶ Almost all RL methods are well-described as GPI.
- ▶ **Push-pull**: Improving the policy will deteriorate value estimation.
- ▶ Well balanced **trade-off between evaluating and improving** is required.




Fig. 3.2: Interpreting generalized policy iteration to switch back and forth between (arbitrary) evaluations and improvement steps

- ▶ DP is applicable for prediction and control problems in MDPs.

Summary

- ▶ DP is applicable for prediction and control problems in MDPs.
- ▶ But requires always full knowledge about the environment (i.e., it is a model-based solution).

Summary

- ▶ DP is applicable for prediction and control problems in MDPs.
- ▶ But requires always full knowledge about the environment (i.e., it is a model-based solution).
- ▶ DP is more efficient than the exhaustive search.

- ▶ DP is applicable for prediction and control problems in MDPs.
- ▶ But requires always full knowledge about the environment (i.e., it is a model-based solution).
- ▶ DP is more efficient than the exhaustive search.
- ▶ But suffers from the curse of dimensionality for large MDPs.

- ▶ DP is applicable for prediction and control problems in MDPs.
- ▶ But requires always full knowledge about the environment (i.e., it is a model-based solution).
- ▶ DP is more efficient than the exhaustive search.
- ▶ But suffers from the curse of dimensionality for large MDPs.
- ▶ (Iterative) policy evaluations and (greedy) improvements solve MDPs.

- ▶ DP is applicable for prediction and control problems in MDPs.
- ▶ But requires always full knowledge about the environment (i.e., it is a model-based solution).
- ▶ DP is more efficient than the exhaustive search.
- ▶ But suffers from the curse of dimensionality for large MDPs.
- ▶ (Iterative) policy evaluations and (greedy) improvements solve MDPs.
- ▶ Both steps can be combined via value iteration.

- ▶ DP is applicable for prediction and control problems in MDPs.
- ▶ But requires always full knowledge about the environment (i.e., it is a model-based solution).
- ▶ DP is more efficient than the exhaustive search.
- ▶ But suffers from the curse of dimensionality for large MDPs.
- ▶ (Iterative) policy evaluations and (greedy) improvements solve MDPs.
- ▶ Both steps can be combined via value iteration.
- ▶ The idea of (generalized) policy iteration is a basic scheme of RL.