

Lecture 04: Monte Carlo Methods

Davit Ghazaryan

February 18, 2025



Table of contents

- 1 General idea and differences to dynamic programming
- 2 Basic Monte Carlo prediction
- 3 Basic Monte Carlo control
- 4 Extensions to Monte Carlo on-policy control
- 5 Monte Carlo off-policy prediction and control

Monte Carlo methods vs. Dynamic Programming

Dynamic programming:

- ▶ **Model-based** prediction and control
- ▶ Planning inside **known MDPs**

Monte Carlo methods vs. Dynamic Programming

Dynamic programming:

- ▶ **Model-based** prediction and control
- ▶ Planning inside **known MDPs**

Monte Carlo methods:

- ▶ **Model-free** prediction and control
- ▶ Estimating value functions and optimize policies in **unknown MDPs**

Monte Carlo methods vs. Dynamic Programming

Dynamic programming:

- ▶ **Model-based** prediction and control
- ▶ Planning inside **known MDPs**

Monte Carlo methods:

- ▶ **Model-free** prediction and control
- ▶ Estimating value functions and optimize policies in **unknown MDPs**
- ▶ But: still assuming finite MDP problems (or problems close to that)

Monte Carlo methods vs. Dynamic Programming

Dynamic programming:

- ▶ **Model-based** prediction and control
- ▶ Planning inside **known MDPs**

Monte Carlo methods:

- ▶ **Model-free** prediction and control
- ▶ Estimating value functions and optimize policies in **unknown MDPs**
- ▶ But: still assuming finite MDP problems (or problems close to that)
- ▶ In general: broad class of computational algorithms relying on **repeated random sampling** to obtain numerical results

Monte-Carlo Reinforcement Learning

- ▶ MC methods learn directly from episodes of experience
- ▶ MC is model-free: no knowledge of MDP transitions / rewards
- ▶ MC learns from complete episodes: no bootstrapping
- ▶ MC uses the simplest possible idea: value = mean return
- ▶ Caveat: can only apply MC to episodic MDPs
 - ▶ All episodes must terminate

General Monte Carlo (MC) methods' characteristics

- ▶ **Learning from experience**, i.e., sequences of samples $\langle s, a, R_{t+1} \rangle$



Fig. 4.1: Monte Carlo port

General Monte Carlo (MC) methods' characteristics

- ▶ **Learning from experience**, i.e., sequences of samples $\langle s, a, R_{t+1} \rangle$
- ▶ Main concept: Estimation by **averaging sample returns**



Fig. 4.1: Monte Carlo port

General Monte Carlo (MC) methods' characteristics

- ▶ **Learning from experience**, i.e., sequences of samples $\langle s, a, R_{t+1} \rangle$
- ▶ Main concept: Estimation by **averaging sample returns**
- ▶ To guarantee well-defined returns: **limited to episodic tasks**



Fig. 4.1: Monte Carlo port

General Monte Carlo (MC) methods' characteristics

- ▶ **Learning from experience**, i.e., sequences of samples $\langle s, a, R_{t+1} \rangle$
- ▶ Main concept: Estimation by **averaging sample returns**
- ▶ To guarantee well-defined returns: **limited to episodic tasks**
- ▶ Consequence: Estimation and policy updates are only possible in an episode-by-episode way compared to step-by-step (online)



Fig. 4.1: Monte Carlo port

Example

- ▶ Scene from the movie "Next"

Table of contents

- 1 General idea and differences to dynamic programming
- 2 Basic Monte Carlo prediction**
- 3 Basic Monte Carlo control
- 4 Extensions to Monte Carlo on-policy control
- 5 Monte Carlo off-policy prediction and control

Task description and basic solution

MC prediction problem statement

- ▶ Estimate state value $v_{\pi}(s)$ for a given policy π .
- ▶ Available are samples $\langle s_{t,j}, a_{t,j}, R_{t+1,j} \rangle$ for episodes $j = 1, \dots, J$.

Task description and basic solution

MC prediction problem statement

- ▶ Estimate state value $v_\pi(s)$ for a given policy π .
- ▶ Available are samples $\langle s_{t,j}, a_{t,j}, R_{t+1,j} \rangle$ for episodes $j = 1, \dots, J$.

MC solution approach:

- ▶ Average returns after visiting state s over episodes $j = 1, \dots$

$$v_\pi(s) \approx \hat{v}_\pi(s) = \frac{1}{J} \sum_{j=1}^J G_{t,j} = \frac{1}{J} \sum_{j=1}^J \sum_{i=0}^{T_j} \gamma^i R_{t+i+1,j}. \quad (4.1)$$

- ▶ Above, T_j denotes the **terminating time step** of each episode j .

Task description and basic solution

MC prediction problem statement

- ▶ Estimate state value $v_\pi(s)$ for a given policy π .
- ▶ Available are samples $\langle s_{t,j}, a_{t,j}, R_{t+1,j} \rangle$ for episodes $j = 1, \dots, J$.

MC solution approach:

- ▶ Average returns after visiting state s over episodes $j = 1, \dots$

$$v_\pi(s) \approx \hat{v}_\pi(s) = \frac{1}{J} \sum_{j=1}^J G_{t,j} = \frac{1}{J} \sum_{j=1}^J \sum_{i=0}^{T_j} \gamma^i R_{t+i+1,j}. \quad (4.1)$$

- ▶ Above, T_j denotes the **terminating time step** of each episode j .
- ▶ **First-visit MC**: Apply (4.1) only to the first state visit per episode.

Task description and basic solution

MC prediction problem statement

- ▶ Estimate state value $v_\pi(s)$ for a given policy π .
- ▶ Available are samples $\langle s_{t,j}, a_{t,j}, R_{t+1,j} \rangle$ for episodes $j = 1, \dots, J$.

MC solution approach:

- ▶ Average returns after visiting state s over episodes $j = 1, \dots$

$$v_\pi(s) \approx \hat{v}_\pi(s) = \frac{1}{J} \sum_{j=1}^J G_{t,j} = \frac{1}{J} \sum_{j=1}^J \sum_{i=0}^{T_j} \gamma^i R_{t+i+1,j}. \quad (4.1)$$

- ▶ Above, T_j denotes the **terminating time step** of each episode j .
- ▶ **First-visit MC**: Apply (4.1) only to the first state visit per episode.
- ▶ **Every-visit MC**: Apply (4.1) each time visiting a certain state per episode (if a state is visited more than one time per episode).

First-Visit Monte-Carlo Policy Evaluation

- ▶ To evaluate state s
- ▶ The **first** time-step t that state s is visited in an episode,
- ▶ Increment counter $N(s) \leftarrow N(s) + 1$
- ▶ Increment total return $S(s) \leftarrow S(s) + G_t$
- ▶ Value is estimated by mean return $\hat{v}_\pi = S(s)/N(s)$
- ▶ By law of large numbers, $\hat{v}_\pi \rightarrow v_\pi(s)$ as $N(s) \rightarrow \infty$

Every-Visit Monte-Carlo Policy Evaluation

- ▶ To evaluate state s
- ▶ **Every** time-step t that state s is visited in an episode,
- ▶ Increment counter $N(s) \leftarrow N(s) + 1$
- ▶ Increment total return $S(s) \leftarrow S(s) + G_t$
- ▶ Value is estimated by mean return $\hat{v}_\pi = S(s)/N(s)$
- ▶ By law of large numbers, $\hat{v}_\pi \rightarrow v_\pi(s)$ as $N(s) \rightarrow \infty$

Algorithmic implementation: MC-based prediction

input: a policy π to be evaluated

output: estimate of v_S^π (i.e., value estimate for all states $s \in \mathcal{S}$)

init: $\hat{v}(s) \forall s \in \mathcal{S}$ arbitrary except $v_0(s) = 0$ if s is terminal

$l(s) \leftarrow$ an empty list for every $s \in \mathcal{S}$

for $j = 1, \dots, J$ *episodes* **do**

Generate an episode following π : $s_0, a_0, R_1, \dots, s_{T_j}, a_{T_j}, R_{T_j+1}$;

Set $G \leftarrow 0$;

for $t = T_j - 1, T_j - 2, T_j - 3, \dots, 0$ *time steps* **do**

$G \leftarrow \gamma G + R_{t+1}$;

if $s_t \notin \langle s_0, s_1, \dots, s_{t-1} \rangle$ **then**

Append G to list $l(s_t)$;

$\hat{v}(s_t) \leftarrow \text{average}(l(s_t))$;

Algo. 4.1: MC state-value prediction (first visit)

Incremental implementation

- ▶ Algo. 4.1 is inefficient due to large memory demand.
- ▶ Better: use **incremental / recursive implementation**.

Incremental implementation

- ▶ Algo. 4.1 is inefficient due to large memory demand.
- ▶ Better: use **incremental / recursive implementation**.
- ▶ The sample mean μ_1, μ_2, \dots of an arbitrary sequence G_1, G_2, \dots is:

$$\begin{aligned}\mu_J &= \frac{1}{J} \sum_{i=1}^J G_i = \frac{1}{J} \left[G_J + \sum_{i=1}^{J-1} G_i \right] \\ &= \frac{1}{J} [G_J + (J-1)\mu_{J-1}] = \mu_{J-1} + \frac{1}{J} [G_J - \mu_{J-1}].\end{aligned}\tag{4.2}$$

Incremental implementation

- ▶ Algo. 4.1 is inefficient due to large memory demand.
- ▶ Better: use **incremental / recursive implementation**.
- ▶ The sample mean μ_1, μ_2, \dots of an arbitrary sequence G_1, G_2, \dots is:

$$\begin{aligned}\mu_J &= \frac{1}{J} \sum_{i=1}^J G_i = \frac{1}{J} \left[G_J + \sum_{i=1}^{J-1} G_i \right] \\ &= \frac{1}{J} [G_J + (J-1)\mu_{J-1}] = \mu_{J-1} + \frac{1}{J} [G_J - \mu_{J-1}].\end{aligned}\tag{4.2}$$

- ▶ If a given decision problem is **non-stationary**, using a forgetting factor $\alpha \in \{\mathbb{R} | 0 < \alpha < 1\}$ allows for dynamic adaption:

$$\mu_J = \mu_{J-1} + \alpha [G_J - \mu_{J-1}].\tag{4.3}$$

Statistical properties of MC-based prediction (1)

First-time visit MC:

- ▶ Each return sample G_J is independent from the others since they were drawn from separate episodes.
- ▶ One receives **i.i.d. data** to estimate $\mathbb{E}[\hat{v}_\pi]$ and consequently this **is bias-free**.
- ▶ The estimate's variance $\text{Var}[\hat{v}_\pi]$ drops with $1/n$ (n : available samples).

Statistical properties of MC-based prediction (1)

First-time visit MC:

- ▶ Each return sample G_J is independent from the others since they were drawn from separate episodes.
- ▶ One receives **i.i.d. data** to estimate $\mathbb{E}[\hat{v}_\pi]$ and consequently this **is bias-free**.
- ▶ The estimate's variance $\text{Var}[\hat{v}_\pi]$ drops with $1/n$ (n : available samples).

Every-time visit MC:

- ▶ Each return sample G_J is not independent from the others since they might be obtained from same episodes.
- ▶ One receives **non-i.i.d.** data to estimate $\mathbb{E}[\hat{v}_\pi]$ and consequently this **is biased** for any $n < \infty$.
- ▶ Only in the limit $n \rightarrow \infty$ one receives $(v_\pi(s) - \mathbb{E}[\hat{v}_\pi(s)]) \rightarrow 0$.

Statistical properties of MC-based prediction (2)

- ▶ State-value estimates for each state are independent.
- ▶ One estimate does not rely on the estimate of other states (no **bootstrapping** compared to DP).
- ▶ Makes MC particularly attractive when one requires state-value knowledge of only one or few states.
 - ▶ Hence, generating episodes starting from the state of interest.




Fig. 4.2: Back-up diagrams for DP (left) and MC (right) prediction: shallow one-step back-ups with bootstrapping vs. deep back-ups over full episodes

MC-based prediction example: forest tree MDP (1)

Let's reuse the forest tree MDP example with *fifty-fifty policy* and discount factor $\gamma = 0.8$ plus disaster probability $\alpha = 0.2$:




Fig. 4.3: Forest MDP with fifty-fifty-policy including state values

MC-based prediction example: forest tree MDP (2)




Fig. 4.4: State-value estimate of forest tree MDP initial state using MC-based prediction over the number of episodes being evaluated (mean and standard deviation are calculated based on 2000 independent runs)

MC estimation of action values

Is a **model available** (i.e., tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$)?

- ▶ **Yes**: state values plus one-step prediction deliver optimal policy.
- ▶ **No**: action values are very useful to directly obtain optimal choices.
- ▶ Recap policy improvement from last lecture.

MC estimation of action values

Is a **model available** (i.e., tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$)?

- ▶ **Yes**: state values plus one-step prediction deliver optimal policy.
- ▶ **No**: action values are very useful to directly obtain optimal choices.
- ▶ Recap policy improvement from last lecture.

Estimating $q_\pi(s, a)$ using MC approach:

- ▶ Analog to state values summarized in Algo. 4.1.
- ▶ Only small extension: a visit refers to a state-action pair (s, a) .
- ▶ First-visit and every-visit variants exist.

MC estimation of action values

Is a **model available** (i.e., tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$)?

- ▶ **Yes**: state values plus one-step prediction deliver optimal policy.
- ▶ **No**: action values are very useful to directly obtain optimal choices.
- ▶ Recap policy improvement from last lecture.

Estimating $q_\pi(s, a)$ using MC approach:

- ▶ Analog to state values summarized in Algo. 4.1.
- ▶ Only small extension: a visit refers to a state-action pair (s, a) .
- ▶ First-visit and every-visit variants exist.

Possible problem when following a deterministic policy π :

- ▶ Certain state-action pairs (s, a) are never visited.
- ▶ Missing degree of exploration.

MC estimation of action values

Is a **model available** (i.e., tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$)?

- ▶ **Yes**: state values plus one-step prediction deliver optimal policy.
- ▶ **No**: action values are very useful to directly obtain optimal choices.
- ▶ Recap policy improvement from last lecture.

Estimating $q_\pi(s, a)$ using MC approach:

- ▶ Analog to state values summarized in Algo. 4.1.
- ▶ Only small extension: a visit refers to a state-action pair (s, a) .
- ▶ First-visit and every-visit variants exist.

Possible problem when following a deterministic policy π :

- ▶ Certain state-action pairs (s, a) are never visited.
- ▶ Missing degree of exploration.
- ▶ Workaround: **exploring starts**, i.e., starting episodes in random state-action pairs (s, a) and thereafter following π .

Table of contents

- 1 General idea and differences to dynamic programming
- 2 Basic Monte Carlo prediction
- 3 Basic Monte Carlo control**
- 4 Extensions to Monte Carlo on-policy control
- 5 Monte Carlo off-policy prediction and control

Applying generalized policy iteration (GPI) to MC control

GPI concept is directly applied to MC framework using action values:

$$\pi_0 \rightarrow \hat{q}_{\pi_0} \rightarrow \pi_1 \rightarrow \hat{q}_{\pi_1} \rightarrow \cdots \pi^* \rightarrow \hat{q}_{\pi^*} . \quad (4.4)$$




Fig. 4.5: Transferring GPI to MC-based control (source: R. Sutton and G. Barto, Reinforcement learning: an introduction, 2018)

Applying generalized policy iteration (GPI) to MC control

GPI concept is directly applied to MC framework using action values:

$$\pi_0 \rightarrow \hat{q}_{\pi_0} \rightarrow \pi_1 \rightarrow \hat{q}_{\pi_1} \rightarrow \cdots \pi^* \rightarrow \hat{q}_{\pi^*} . \quad (4.4)$$

- Degree of freedom: Choose number of episodes to approximate \hat{q}_{π_i} .




Fig. 4.5: Transferring GPI to MC-based control (source: R. Sutton and G. Barto, Reinforcement learning: an introduction, 2018)

Applying generalized policy iteration (GPI) to MC control

GPI concept is directly applied to MC framework using action values:

$$\pi_0 \rightarrow \hat{q}_{\pi_0} \rightarrow \pi_1 \rightarrow \hat{q}_{\pi_1} \rightarrow \cdots \pi^* \rightarrow \hat{q}_{\pi^*}. \quad (4.4)$$

- ▶ Degree of freedom: Choose number of episodes to approximate \hat{q}_{π_i} .
- ▶ Policy improvement is done by greedy choices:

$$\pi(s) = \arg \max_a q(s, a) \quad \forall s \in \mathcal{S}. \quad (4.5)$$




Fig. 4.5: Transferring GPI to MC-based control (source: R. Sutton and G. Barto, Reinforcement learning: an introduction, 2018)

Policy improvement theorem

Assuming that one is operating in an **unknown MDP**, the policy improvement theorem is still valid for MC-based control:

Policy improvement for MC-based control

$$\begin{aligned} q_{\pi_{i+1}}(s, \pi_{i+1}(s)) &= q_{\pi_i}(s, \arg \max_a q_{\pi_i}(s, a)), \\ &= \max_a q_{\pi_i}(s, a), \\ &\geq q_{\pi_i}(s, \pi_i(s)), \\ &\geq v_{\pi_i}(s). \end{aligned} \tag{4.6}$$

- ▶ Each π_{i+1} is uniformly better or just as good (if optimal) as π_i .
- ▶ Assumption: All state-action pairs are evaluated due to sufficient exploration.
 - ▶ For example using exploring starts.

Algorithmic implementation: MC-based control

output: Optimal deterministic policy π^*

init: $\pi_{i=0}(s) \in \mathcal{A}$ arbitrarily $\forall s \in \mathcal{S}$

$\hat{q}(s, a)$ arbitrarily $\forall \{s \in \mathcal{S}, a \in \mathcal{A}\}$

$n(s, a) \leftarrow$ an empty list for state-action visits $\forall \{s \in \mathcal{S}, a \in \mathcal{A}\}$

repeat

$i \leftarrow i + 1$;

Choose $\{s_0, a_0\}$ randomly such that all pairs have probability > 0 ;

Generate an episode from $\{s_0, a_0\}$ following π_i until termination step T_i ;

Set $G \leftarrow 0$;

for $k = T_i - 1, T_i - 2, T_i - 3, \dots, 0$ *time steps* **do**

$G \leftarrow \gamma G + R_{t+1}$;

if $\{s_t, a_t\} \notin \langle \{s_0, a_0\}, \dots, \{s_{t-1}, a_{t-1}\} \rangle$ **then**

$n(s_t, a_t) \leftarrow n(s_t, a_t) + 1$;

$\hat{q}(s_t, a_t) \leftarrow \hat{q}(s_t, a_t) + 1/n(s_t, a_t) \cdot (G - \hat{q}(s_t, a_t))$;

$\pi_i(s_t) \leftarrow \arg \max_a \hat{q}(s_t, a)$;

until $\pi_{i+1} = \pi_i$;

Algo. 4.2: MC-based control using exploring starts (first visit)

Table of contents

- 1 General idea and differences to dynamic programming
- 2 Basic Monte Carlo prediction
- 3 Basic Monte Carlo control
- 4 Extensions to Monte Carlo on-policy control**
- 5 Monte Carlo off-policy prediction and control

▶ On-policy learning

- ▶ Evaluate or improve the policy used to make decisions.
- ▶ Agent picks actions based on its own policy.
- ▶ Exploring starts (ES) is an on-policy method example.
- ▶ However: ES is a restrictive assumption and not always applicable (in some cases the starting state-action pair cannot be chosen freely).

Off- and on-policy learning

▶ On-policy learning

- ▶ Evaluate or improve the policy used to make decisions.
- ▶ Agent picks actions based on its own policy.
- ▶ Exploring starts (ES) is an on-policy method example.
- ▶ However: ES is a restrictive assumption and not always applicable (in some cases the starting state-action pair cannot be chosen freely).

▶ Off-policy learning

- ▶ Evaluate or improve a policy different from that used to generate data.
- ▶ Agent cannot apply own actions.
- ▶ Will be focused in the next sections.

ε -greedy as an on-policy alternative

- ▶ Exploration requirement:
 - ▶ Visit all state-action pairs with probability:

$$\pi(a|s) > 0 \quad \forall \{s \in \mathcal{S}, a \in \mathcal{A}\} . \quad (4.7)$$

- ▶ Policies with this characteristic are called: **soft**.
 - ▶ Level of exploration can be tuned during the learning process.

ε -greedy as an on-policy alternative

- ▶ Exploration requirement:

- ▶ Visit all state-action pairs with probability:

$$\pi(a|s) > 0 \quad \forall \{s \in \mathcal{S}, a \in \mathcal{A}\} . \quad (4.7)$$

- ▶ Policies with this characteristic are called: **soft**.
 - ▶ Level of exploration can be tuned during the learning process.

- ▶ **ε -greedy on-policy learning**

- ▶ With probability ε the agent's choice (i.e., the policy output) is overwritten with a random action.
 - ▶ Probability of all non-greedy actions:

$$\varepsilon/|\mathcal{A}| . \quad (4.8)$$

- ▶ Probability of the greedy action:

$$1 - \varepsilon + \varepsilon/|\mathcal{A}| . \quad (4.9)$$

- ▶ Above, $|\mathcal{A}|$ is the size of the action space.

Algorithmic implementation ε -greedy MC-control

output: Optimal ε -greedy policy $\pi^*(a|s)$, **parameter:** $\varepsilon \in \{\mathbb{R} | 0 < \varepsilon \ll 1\}$
init: $\pi_{i=0}(a|s)$ arbitrarily soft $\forall \{s \in \mathcal{S}, a \in \mathcal{A}\}$
 $\hat{q}(s, a)$ arbitrarily $\forall \{s \in \mathcal{S}, a \in \mathcal{A}\}$
 $n(s, a) \leftarrow$ an empty list counting state-action visits $\forall \{s \in \mathcal{S}, a \in \mathcal{A}\}$
repeat
 Generate an episode following π_i : $s_0, a_0, R_1, \dots, s_{T_j}, a_{T_j}, R_{T_j+1}$;
 $i \leftarrow i + 1$;
 Set $G \leftarrow 0$;
 for $t = T_i - 1, T_i - 2, T_i - 3, \dots, 0$ *time steps* **do**
 $G \leftarrow \gamma G + R_{t+1}$;
 if $\{s_t, a_t\} \notin \langle \{s_0, a_0\}, \dots, \{s_{t-1}, a_{t-1}\} \rangle$ **then**
 $n(s_t, a_t) \leftarrow n(s_t, a_t) + 1$;
 $\hat{q}(s_t, a_t) \leftarrow \hat{q}(s_t, a_t) + 1/n(s_t, a_t) \cdot (G - \hat{q}(s_t, a_t))$;
 $\tilde{a} \leftarrow \arg \max_a \hat{q}(s_t, a)$;
 $\pi_i(a|s_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}|, & a = \tilde{a} \\ \varepsilon/|\mathcal{A}|, & a \neq \tilde{a} \end{cases}$;
 until $\pi_{i+1} = \pi_i$;

Algo. 4.3: MC-based control using ε -greedy approach

ε -greedy policy improvement (1)

Theorem 4.1: Policy improvement for ε -greedy policy

Given an MDP, for any ε -greedy policy π the ε -greedy policy π' with respect to q_π is an improvement, i.e., $v_{\pi'} > v_\pi \quad \forall s \in \mathcal{S}$.

ε -greedy policy improvement (1)

Theorem 4.1: Policy improvement for ε -greedy policy

Given an MDP, for any ε -greedy policy π the ε -greedy policy π' with respect to q_π is an improvement, i.e., $v_{\pi'} > v_\pi \quad \forall s \in \mathcal{S}$.

Small proof:

$$\begin{aligned} q_\pi(s, \pi'(s)) &= \sum_a \pi'(a|s) q_\pi(s, a), \\ &= \frac{\varepsilon}{|\mathcal{A}|} \sum_a q_\pi(s, a) + (1 - \varepsilon) \max_a q_\pi(s, a), \\ &\geq \frac{\varepsilon}{|\mathcal{A}|} \sum_a q_\pi(s, a) + (1 - \varepsilon) \sum_a \frac{\pi(a|s) - \frac{\varepsilon}{|\mathcal{A}|}}{1 - \varepsilon} q_\pi(s, a). \end{aligned} \tag{4.10}$$

In the inequality line, the second term is the weighted sum over action values given an ε -greedy policy. This weighted sum will always be smaller than or equal to $\max_a q_\pi(s, a)$.

Continuation:

$$\begin{aligned} q_{\pi}(s, \pi'(s)) &\geq \frac{\varepsilon}{|\mathcal{A}|} \sum_a q_{\pi}(s, a) + (1 - \varepsilon) \sum_a \frac{\pi(a|s) - \frac{\varepsilon}{|\mathcal{A}|}}{1 - \varepsilon} q_{\pi}(s, a), \\ &= \frac{\varepsilon}{|\mathcal{A}|} \sum_a (q_{\pi}(s, a) - q_{\pi}(s, a)) + \sum_a \pi(a|s) q_{\pi}(s, a), \\ &= \sum_a \pi(a|s) q_{\pi}(s, a), \\ &= v_{\pi}(s). \end{aligned} \tag{4.11}$$

Continuation:

$$\begin{aligned} q_{\pi}(s, \pi'(s)) &\geq \frac{\varepsilon}{|\mathcal{A}|} \sum_a q_{\pi}(s, a) + (1 - \varepsilon) \sum_a \frac{\pi(a|s) - \frac{\varepsilon}{|\mathcal{A}|}}{1 - \varepsilon} q_{\pi}(s, a), \\ &= \frac{\varepsilon}{|\mathcal{A}|} \sum_a (q_{\pi}(s, a) - q_{\pi}(s, a)) + \sum_a \pi(a|s) q_{\pi}(s, a), \\ &= \sum_a \pi(a|s) q_{\pi}(s, a), \\ &= v_{\pi}(s). \end{aligned} \tag{4.11}$$

- ▶ Policy improvement theorem is still valid when comparing ε -greedy policies against each other.
- ▶ But: There might be a non- ε -greedy policy which is better.

MC-based control example: forest tree MDP (1)




Fig. 4.6: Different estimates of forest tree MDP ($\alpha = 0.2, \gamma = 0.8$) using MC control with $\varepsilon = 0.2$ over the number of episodes. Mean is red and the standard deviation is light blue, both calculated based on 2000 independent runs.

MC-based control example: forest tree MDP (2)




Fig. 4.7: Different estimates of forest tree MDP ($\alpha = 0.2, \gamma = 0.8$) using MC control with $\varepsilon = 0.05$ over the number of episodes. Mean is red and the standard deviation is light blue, both calculated based on 2000 independent runs.

MC-based control example: forest tree MDP (3)

Observations on forest tree MDP with ϵ -greedy MC-based control:

- ▶ Rather slow convergence rate: quite a number of episodes is required.

MC-based control example: forest tree MDP (3)

Observations on forest tree MDP with ε -greedy MC-based control:

- ▶ Rather slow convergence rate: quite a number of episodes is required.
- ▶ Significant uncertainty present in a single sequence.

MC-based control example: forest tree MDP (3)

Observations on forest tree MDP with ϵ -greedy MC-based control:

- ▶ Rather slow convergence rate: quite a number of episodes is required.
- ▶ Significant uncertainty present in a single sequence.
- ▶ Later states are less often visited and, therefore, more uncertain.

MC-based control example: forest tree MDP (3)

Observations on forest tree MDP with ε -greedy MC-based control:

- ▶ Rather slow convergence rate: quite a number of episodes is required.
- ▶ Significant uncertainty present in a single sequence.
- ▶ Later states are less often visited and, therefore, more uncertain.
- ▶ Exploration is controlled by ε : in a totally greedy policy the state $s = 3$ is not visited at all. With ε -greedy this state is visited occasionally.

MC-based control example: forest tree MDP (3)

Observations on forest tree MDP with ε -greedy MC-based control:

- ▶ Rather slow convergence rate: quite a number of episodes is required.
- ▶ Significant uncertainty present in a single sequence.
- ▶ Later states are less often visited and, therefore, more uncertain.
- ▶ Exploration is controlled by ε : in a totally greedy policy the state $s = 3$ is not visited at all. With ε -greedy this state is visited occasionally.
- ▶ Nevertheless, the above figures highlight that MC-based control for the forest tree MDP tend towards the optimal policies discovered by dynamic programming.

Greedy in the Limit with Infinite Exploration (GLIE)

Definition 4.1: Greedy in the limit with infinite exploration (GLIE)

A learning policy π is called GLIE if it satisfies the following two properties:

- ▶ If a state is visited infinitely often, then each action is chosen infinitely often:

$$\lim_{i \rightarrow \infty} \pi_i(a|s) = 1 \quad \forall \{s \in \mathcal{S}, a \in \mathcal{A}\} . \quad (4.12)$$

- ▶ In the limit ($i \rightarrow \infty$) the learning policy is greedy with respect to the learned action value:

$$\lim_{i \rightarrow \infty} \pi_i(a|s) = \pi(s) = \arg \max_a q(s, a) \quad \forall s \in \mathcal{S} . \quad (4.13)$$

Theorem 4.2: Optimal decision using MC-control with ε -greedy

MC-based control using ε -greedy exploration is GLIE, if ε is decreased at rate

$$\varepsilon_i = \frac{1}{i} \quad (4.14)$$

with i being the increasing episode index. In this case,

$$\hat{q}(s, a) = q^*(s, a) \quad (4.15)$$

follows.

Theorem 4.2: Optimal decision using MC-control with ε -greedy

MC-based control using ε -greedy exploration is GLIE, if ε is decreased at rate

$$\varepsilon_i = \frac{1}{i} \quad (4.14)$$

with i being the increasing episode index. In this case,

$$\hat{q}(s, a) = q^*(s, a) \quad (4.15)$$

follows.

Remarks:

- ▶ Limited feasibility: infinite number of episodes required.
- ▶ ε -greedy is an undirected and unmonitored random exploration strategy. Can that be the most efficient way of learning?

Table of contents

- 1 General idea and differences to dynamic programming
- 2 Basic Monte Carlo prediction
- 3 Basic Monte Carlo control
- 4 Extensions to Monte Carlo on-policy control
- 5 Monte Carlo off-policy prediction and control

Off-policy learning background

Drawback of on-policy learning:

- ▶ Only a compromise: comes with inherent exploration but at the cost of learning action values for a **near-optimal policy**.

Off-policy learning background

Drawback of on-policy learning:

- ▶ Only a compromise: comes with inherent exploration but at the cost of learning action values for a **near-optimal policy**.

Idea off-policy learning:

- ▶ Use two separated policies:
 - ▶ **Behavior policy** $b(a|s)$: explores in order to generate experience.
 - ▶ **Target policy** $\pi(a|s)$: learns from that experience to become the optimal policy.

Off-policy learning background

Drawback of on-policy learning:

- ▶ Only a compromise: comes with inherent exploration but at the cost of learning action values for a **near-optimal policy**.

Idea off-policy learning:

- ▶ Use two separated policies:
 - ▶ **Behavior policy** $b(a|s)$: explores in order to generate experience.
 - ▶ **Target policy** $\pi(a|s)$: learns from that experience to become the optimal policy.
- ▶ Use cases:
 - ▶ Learn from observing humans or other agents/controllers.

Off-policy learning background

Drawback of on-policy learning:

- ▶ Only a compromise: comes with inherent exploration but at the cost of learning action values for a **near-optimal policy**.

Idea off-policy learning:

- ▶ Use two separated policies:
 - ▶ **Behavior policy** $b(a|s)$: explores in order to generate experience.
 - ▶ **Target policy** $\pi(a|s)$: learns from that experience to become the optimal policy.
- ▶ Use cases:
 - ▶ Learn from observing humans or other agents/controllers.
 - ▶ Re-use experience generated from old policies (π_0, π_1, \dots) .

Off-policy learning background

Drawback of on-policy learning:

- ▶ Only a compromise: comes with inherent exploration but at the cost of learning action values for a **near-optimal policy**.

Idea off-policy learning:

- ▶ Use two separated policies:
 - ▶ **Behavior policy** $b(a|s)$: explores in order to generate experience.
 - ▶ **Target policy** $\pi(a|s)$: learns from that experience to become the optimal policy.
- ▶ Use cases:
 - ▶ Learn from observing humans or other agents/controllers.
 - ▶ Re-use experience generated from old policies (π_0, π_1, \dots) .
 - ▶ Learn about multiple policies while following one policy.

Off-policy prediction problem statement

MC off-policy prediction problem statement

- ▶ Estimate v_π and/or q_π while following $b(a|s)$.
- ▶ Both policies are considered fixed (prediction assumption).

Off-policy prediction problem statement

MC off-policy prediction problem statement

- ▶ Estimate v_π and/or q_π while following $b(a|s)$.
- ▶ Both policies are considered fixed (prediction assumption).

Requirement:

- ▶ **Coverage:** Every action taken under π must be (at least occasionally) taken under b , too.
Hence, it follows:

$$\pi(a|s) > 0 \Rightarrow b(a|s) > 0 \quad \forall \{s \in \mathcal{S}, a \in \mathcal{A}\}. \quad (4.16)$$

Off-policy prediction problem statement

MC off-policy prediction problem statement

- ▶ Estimate v_π and/or q_π while following $b(a|s)$.
- ▶ Both policies are considered fixed (prediction assumption).

Requirement:

- ▶ **Coverage:** Every action taken under π must be (at least occasionally) taken under b , too.
Hence, it follows:

$$\pi(a|s) > 0 \Rightarrow b(a|s) > 0 \quad \forall \{s \in \mathcal{S}, a \in \mathcal{A}\}. \quad (4.16)$$

- ▶ Consequences from that:
 - ▶ In any state b is not identical to π , b must be stochastic.
 - ▶ Nevertheless: π might be deterministic (e.g., control applications) or stochastic.

Importance sampling

Probability of observing a certain trajectory on random variables $A_t, S_{t+1}, A_{t+1}, \dots, S_T$ starting in S_t while following π :

$$\begin{aligned}\mathbb{P}[A_t, S_{t+1}, A_{t+1}, \dots, S_T | A_t, \pi] &= \pi(A_t | S_t) p(S_{t+1} | S_t, A_t) \pi(A_{t+1} | S_{t+1}) \cdots, \\ &= \prod_t^{T-1} \pi(A_t | S_t) p(S_{t+1} | S_t, A_t).\end{aligned}\tag{4.17}$$

Above p is the state-transition probability.

Importance sampling

Probability of observing a certain trajectory on random variables $A_t, S_{t+1}, A_{t+1}, \dots, S_T$ starting in S_t while following π :

$$\begin{aligned}\mathbb{P}[A_t, S_{t+1}, A_{t+1}, \dots, S_T | A_t, \pi] &= \pi(A_t | S_t) p(S_{t+1} | S_t, A_t) \pi(A_{t+1} | S_{t+1}) \cdots, \\ &= \prod_t^{T-1} \pi(A_t | S_t) p(S_{t+1} | S_t, A_t).\end{aligned}\tag{4.17}$$

Above p is the state-transition probability.

Definition 4.2: Importance sampling ratio

The relative probability of a trajectory under the target and behavior policy, the importance sampling ratio, from sample step k to T is:

$$\rho_{k:T} = \frac{\prod_t^{T-1} \pi(A_t | S_t) p(S_{t+1} | S_t, A_t)}{\prod_t^{T-1} b(A_t | S_t) p(S_{t+1} | S_t, A_t)} = \frac{\prod_t^{T-1} \pi(A_t | S_t)}{\prod_t^{T-1} b(A_t | S_t)}.\tag{4.18}$$

Importance sampling for Monte Carlo prediction

Definition 4.3: State-value estimation via Monte Carlo importance sampling

Estimating the state value v_π following a behavior policy b using (ordinary) importance sampling (OIS) results in scaling and averaging the sampled returns by the importance sampling ratio per episode:

$$\hat{v}_\pi(s_t) = \frac{\sum_{t \in \mathcal{T}(s_t)} \rho_{t:T(t)} G_t}{|\mathcal{T}(s_t)|}. \quad (4.19)$$

Notation remark:

- ▶ $\mathcal{T}(s_t)$: set of all time steps in which the state s_t is visited.
- ▶ $T(t)$: Termination of a specific episode starting from t .

Importance sampling for Monte Carlo prediction

Definition 4.3: State-value estimation via Monte Carlo importance sampling

Estimating the state value v_π following a behavior policy b using (ordinary) importance sampling (OIS) results in scaling and averaging the sampled returns by the importance sampling ratio per episode:

$$\hat{v}_\pi(s_t) = \frac{\sum_{t \in \mathcal{T}(s_t)} \rho_{t:T(t)} G_t}{|\mathcal{T}(s_t)|}. \quad (4.19)$$

Notation remark:

- ▶ $\mathcal{T}(s_t)$: set of all time steps in which the state s_t is visited.
- ▶ $T(t)$: Termination of a specific episode starting from t .

General remark:

- ▶ From (4.18) it can be seen that \hat{v} is bias-free (first-visit assumption).
- ▶ However, if ρ is large (distinctly different policies) the estimate's variance is large (i.e., uncertain for small numbers of samples).

Off-policy Monte Carlo control

Just put everything together:

- ▶ MC-based control utilizing GPI (cf. Fig. 4.5),
- ▶ Off-policy learning based on importance sampling (or variants like weighted importance sampling, cf. Barto/Sutton book chapter 5.5).

Off-policy Monte Carlo control

Just put everything together:

- ▶ MC-based control utilizing GPI (cf. Fig. 4.5),
- ▶ Off-policy learning based on importance sampling (or variants like weighted importance sampling, cf. Barto/Sutton book chapter 5.5).

Requirement for off-policy MC-based control:

- ▶ **Coverage**: behavior policy b has nonzero probability of selecting actions that might be taken by the target policy π .
- ▶ **Consequence**: behavior policy b is **soft** (e.g., ϵ -soft).

Summary

- ▶ MC methods allow model-free learning of value functions and optimal policies from experience in the form of sampled episodes.
- ▶ Using deep back-ups over full episodes, MC is largely based on averaging returns.

Summary

- ▶ MC methods allow model-free learning of value functions and optimal policies from experience in the form of sampled episodes.
- ▶ Using deep back-ups over full episodes, MC is largely based on averaging returns.
- ▶ MC-based control reuses generalized policy iteration (GPI), i.e., mixing policy evaluation and improvement.

Summary

- ▶ MC methods allow model-free learning of value functions and optimal policies from experience in the form of sampled episodes.
- ▶ Using deep back-ups over full episodes, MC is largely based on averaging returns.
- ▶ MC-based control reuses generalized policy iteration (GPI), i.e., mixing policy evaluation and improvement.
- ▶ Maintaining sufficient exploration is important:
 - ▶ Exploring starts: not feasible in all applications but simple.
 - ▶ On-policy ϵ -greedy learning: trade-off between optimality and exploration cannot be resolved easily.
 - ▶ Off-policy learning: agent learns about a (possibly deterministic) target policy from an exploratory, soft behavior policy.

Summary

- ▶ MC methods allow model-free learning of value functions and optimal policies from experience in the form of sampled episodes.
- ▶ Using deep back-ups over full episodes, MC is largely based on averaging returns.
- ▶ MC-based control reuses generalized policy iteration (GPI), i.e., mixing policy evaluation and improvement.
- ▶ Maintaining sufficient exploration is important:
 - ▶ Exploring starts: not feasible in all applications but simple.
 - ▶ On-policy ϵ -greedy learning: trade-off between optimality and exploration cannot be resolved easily.
 - ▶ Off-policy learning: agent learns about a (possibly deterministic) target policy from an exploratory, soft behavior policy.
- ▶ Importance sampling transforms expectations from the behavior to the target policy.
 - ▶ This estimation task comes with a bias-variance-dilemma.
 - ▶ Slow learning can result from ineffective experience usage in MC methods.